




# Chapter 6: Synchronization Tools



Operating System Concepts – 10<sup>th</sup> Edition



## Chapter 6: Synchronization Tools

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Mutex Locks
- Semaphores





## Objectives

- Describe the **critical-section problem** and illustrate the **race condition**
- Describe hardware solutions to the critical-section problem using compare-and-swap operations, and atomic variables
- Demonstrate how mutex locks, semaphores, and condition variables can be used to solve the critical section problem



## Background

- Processes can execute concurrently
  - Processes may be interrupted at any time, partially completing execution, due to a variety of reasons.
- Concurrent access to any shared data may result in **data inconsistency**
- Maintaining data consistency requires OS mechanisms to ensure the **orderly execution** of cooperating processes





## Illustration of the Problem

- Think about the **Producer-Consumer** problem
  - An integer **counter** is used to keep track of the number of buffers occupied.
    - Initially, **counter** is set to 0
    - It is **incremented** each time by the **producer** after it produces an item and places in the buffer
    - It is **decremented** each time by the **consumer** after it consumes an item in the buffer.
- counter is a shared variable updated by P & C at same time



## Producer-Consumer Problem

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

Producer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

Consumer





## Race Condition

- `counter++` could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

*Assembly code  
→ executing 3 instructions*

- `counter--` could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with "count = 5" initially:

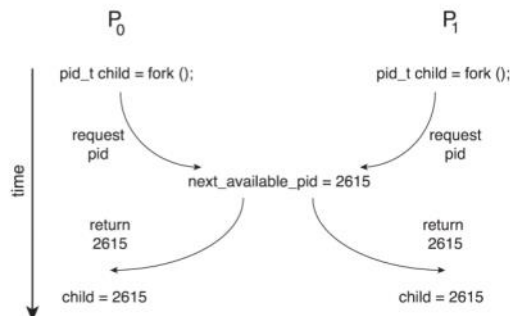
|                      |  |                 |
|----------------------|--|-----------------|
| S0: producer execute | <code>register1 = counter</code>       | {register1 = 5} |
| S1: producer execute | <code>register1 = register1 + 1</code> | {register1 = 6} |
| S2: consumer execute | <code>register2 = counter</code>       | {register2 = 5} |
| S3: consumer execute | <code>register2 = register2 - 1</code> | {register2 = 4} |
| S4: producer execute | <code>counter = register1</code>       | {counter = 6}   |
| S5: consumer execute | <code>counter = register2</code>       | {counter = 4}   |

*can't be sure which  
execution will run  
first.*



## Race Condition

- Processes  $P_0$  and  $P_1$  are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (`pid`)



- Unless there is `mutual exclusion`, the same `pid` could be assigned to two different processes!





## Critical Section Problem

- A **Race Condition** is an undesirable situation where **several processes access or/and manipulate a shared data concurrently** and the outcome of the executions depends on the particular order in which the accesses or executions take place - The results depend on the timing execution of programs. With some bad luck (i.e., context switches that occur at untimely points during execution), the result become **non-deterministic**

↳ for same inputs can have diff outputs

- Consider a system with  $n$  processes  $\{P_0, P_1, \dots, P_{n-1}\}$
- A process has a **Critical Section segment** of code (can be short or long), during which
  - A process or thread may be changing shared variables, updating a table, writing a file, etc.
  - We need to ensure when one process is in Critical Section, no other can be in its critical section
  - In a way, **mutual exclusion** and **critical section** imply the same thing
- Critical section problem is to design a protocol to solve this
  - Specifically, each process must ask permissions before entering a critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

↳ only 1 process runs at a time



## Critical Section

- The general structure of process  $P_i$  is

```
do {  
    entry section  
    critical section ← editing shared variables  
    exit section  
    remainder section  
} while (true);
```





## Solution to Critical-Section Problem

1. **Mutual exclusion** - If process  $P_i$  is executing in its critical section, no other processes can be executing in their critical sections
  2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical sections, the selection of a process that will enter the critical section next *cannot* be postponed *indefinitely* – selection of one process entering
  3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted – any waiting process
- Assume that each process executes at a nonzero speed, and there is no assumption concerning **relative speed** of each individual process

← keep CPU busy  
 ← don't want selection process to take too long since it wastes CPU burst  
 ← fairness  
 ← reduces starvation



## Critical-Section Problem in Kernel

- **Kernel code** - the code the operating system is running, is subject to several possible race conditions

Problem examples:

- A kernel data structure that maintains a list of all open files can be updated by multiple kernel processes, i.e., two processes were to open files simultaneously
- Other kernel data structures such as the one maintaining memory allocation, process lists, interrupt handling etc.

ACB, TCB are maintained in a LL

→ accessing & changing these LL's simultaneously are an example of critical section in the kernel

CSP are most common problems in the kernel

- Two general approaches are used to handle critical sections in operating system, depending on whether the kernel is preemptive or non-preemptive

- **Preemptive** – allows preemption of process when running in the kernel mode, not free from the race condition, and increasingly more difficult in SMP architectures.
- **Non-preemptive** – runs until exiting the kernel mode, blocks, or voluntarily yields CPU. This is essentially free of race conditions in the kernel mode, possibly used in single-processor system

← most OS use preemptive CS

← sacrifices parallelism  
 i.e. when a processor makes I/O request all other processes must wait





## Synchronization Tools

- Many systems provide hardware support for implementing the critical section code. On uniprocessor systems – it could simply disable interrupts, currently running code would execute without being preempted or interrupted. But this is generally inefficient on multiprocessor systems
- Operating systems provide hardware and high level API support for critical section code

| Programs        | Share Programs   |
|-----------------|--|
| Hardware        | Load/Store, Disable Interrupts, Test&Set, Compare&Swap |
| High level APIs | Locks, Semaphores<br>↓<br>Mutex                        |

interrupts cannot preempt



## Synchronization Hardware

- Modern OS provides special **atomic** hardware instructions
  - ▶ **Atomic** = non-interruptible
- There are two commonly used atomic hardware instructions, which can be used to construct more sophisticated synchronization tools
  - Either test a memory word and set a value – `Test_and_Set()`
  - Or swap contents of two memory words – `Compare_and_Swap()`







## test\_and\_set Instruction

### Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

*Handwritten notes:*  
 ↗ get original value  
 ↗ set original value to true  
 ↗ set target to true  
 ↗ return original value of target  
 i.e. if T = false  
 → set T = true  
 → return false



## Solution using test\_and\_set()

### Shared boolean variable **lock**, initialized to **FALSE**

### Solution:

*Handwritten note:* All processes execute this loop

```
do {
    while (test_and_set(&lock))
        ; /* do nothing */
    /* critical section */
    lock = false;
    /* remainder section */
} while (true);
```

*Handwritten notes:*  
 blocking [ ... ]  
 ↗ First process would set Lock to T & return F  
 ↗ first process access this CS, releasing the lock  
 ↗ lock = F  
 ↗ lock is shared variable b/w all 'n' processes







## compare\_and\_swap Instruction

### Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

← always return ORIGINAL VALUE [0v]



## Solution using compare\_and\_swap

- Shared Boolean variable lock initialized to FALSE; Each process has a local Boolean variable key

### Solution:

```
do {
    while (compare_and_swap(&lock, 0, 1) != 0)
        ; /* do nothing */
    /* critical section */
    lock = 0;
    /* remainder section */
} while (true);
```

if  $L = 0$ ,  $L$  is free  
if  $L = 1$ ,  $L$  is acquired by a process

← if return value = 1 then  $L$  is already acquired

← this doesn't ensure bounded waiting

C&S can do more complicated jobs compared to T&S



## Bounded-waiting Mutual Exclusion with test\_and\_set

processes:  $0, 1, \dots, n-1$

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false; /* no one is waiting, so release the lock */
    else
        waiting[j] = false; /* Unblock process j */
    /* remainder section */
}
```

← squared away  
← key tracks the lock status

← neighbour  
← search next process = cs

← at most wait for  $n-1$  times until your turn. But we don't know time

```

else
    waiting[j] = false; /* Unblock process j */
/* remainder section */
} while (true);

```



## Sketch Proof

- **Mutual-exclusion:**  $P_i$  enters its critical section only if either `waiting[i]==false` or `key==false`. The value of `key` can become false only if `test_and_set()` is executed. Only the first process to execute `test_and_set()` will find `key==false`; all others must wait. The variable `waiting[i]` can become false only if another process leaves its critical section; only one `waiting[i]` is set to false, thus maintaining the mutual-exclusion requirement.
- **Progress:** since a process existing its critical section either sets `lock` to false or sets `waiting[j]` to false. Both allow a process that is waiting to enter its critical section to proceed.
- **Bounded-waiting:** when a process leaves its critical section, it scans the array `waiting` in cyclic order  $\{i+1, i+2, \dots, n-1, 0, 1, \dots, i-1\}$ . It designates the first process in this ordering that is in the entry section (`waiting[j]==true`) as the next one to enter the critical section. Any process waiting to enter its critical section will thus do so within  $n-1$  turns.

← selection can't take indefinitely long.



## Atomic Variables

- Typically, instructions such as compare-and-swap are used as building blocks for other (more sophisticated) synchronization tools.
- One tool is an **atomic variable** that provides *atomic* (uninterruptible) updates on basic data types such as integers and Booleans.
- For example, the `increment()` operation on the **atomic variable sequence** ensures **sequence is incremented without interruption - `increment(&sequence)`**;

- The `increment()` function can be implemented as follows: ← busy waiting

```

void increment(atomic_int *v)
{
    int temp;

    do {
        temp = *v;

        while (temp !=
            (compare_and_swap(v, temp, temp+1));
    }

```

← original  
 ↓ expected  
 ↓ temp

← can be used





# Mutex Locks *→ mutually exclusive*

- OS builds software tools to solve the critical section problem
- The simplest tool that most OSes use is **mutex lock**
- To access the critical regions with it by first **acquire()** a lock then **release()** it afterwards
  - Boolean variable indicating if lock is available or not
- Calls to **acquire()** and **release()** must be **atomic** (non-interruptible)
  - Usually implemented via hardware atomic instructions
- But this solution requires **busy waiting**. This lock therefore called a **spinlock**
  - **Spinlock** wastes CPU cycles due to busy waiting, but it has one advantage in that no context switch is required when a process must wait on a lock, and a context switch may take considerable time. Thus, when locks are expected to be held for short times, spinlock is useful
  - Spinlocks are often used in multiprocessor systems where one thread can "spin" on one processor while another thread performs its critical section on another processor



Operating System Concepts – 10th Edition

6.22



## acquire() and release()

*Spin lock logic:*

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
  
release() {  
    available = true;  
}  
  
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

Solutions based on the idea of **lock** to protect critical section

- Operations are **atomic** (non-interruptible) – at most one thread acquires a lock at a time
- Lock before entering critical section for accessing share data
- Unlock upon departure from critical section after accessing shared data
- Wait if locked - all synchronization involves busy waiting, should "sleep" or "block" if waiting for a long time

*← not real code just example of the logic*



Operating System Concepts – 10th Edition

6.23



## Semaphore

- Semaphore **S** – non-negative integer variable, can be considered as a generalized lock
  - First defined by Dijkstra in late 1960s. It can behave similarly as mutex lock, but it has more sophisticated usage - the main synchronization primitive used in original UNIX
- Two standard operations modify **S**: **wait()** and **signal()**
  - Originally called **P()** and **V()**, where **P()** stands for “**proberen**” (to test) and **V()** stands for “**verhogen**” (to increment) in Dutch
- It is essential that semaphore operations are executed **atomically**, which guarantees that no more than one process can execute **wait()** and **signal()** operations on the same semaphore at the same time.
- The semaphore can only be accessed via these two atomic operations except initialization

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
signal (S) {  
    S++;  
}
```



## Semaphore Usage

- **Counting semaphore** – An integer value can range over an unrestricted domain
  - Counting semaphore can be used to control access to a given set of resources consisting of a finite number of instances; semaphore value is initialized to the number of resources available
- **Binary semaphore** – integer value can range only between 0 and 1
  - This can behave like **mutex locks**, can also be used in different ways
- This can also be used to solve various synchronization problems
- Consider **P<sub>1</sub>** and **P<sub>2</sub>** that shares a common semaphore **synch**, initialized to 0; it ensures that **P<sub>1</sub>** process executes **S<sub>1</sub>** before **P<sub>2</sub>** process executes **S<sub>2</sub>**

```
P1:  
    S1;  
    signal(synch);  
P2:  
    wait(synch);  
    S2;
```

★ How many semaphores needed for  $n$  processes  
→  $\frac{n}{2}$  or  $(n-1)$



← keep track of available instances  
in a system i.e. empty buffer slot  
OR can be used as mutex lock



## Semaphore Implementation with no Busy waiting

- Each semaphore is associated with a **waiting queue**
- Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record on the queue *← linked list*
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue
  - **wakeup** – remove one of processes in the waiting queue and place it on the ready queue
- Semaphore values may become negative, whereas this value can never be negative under the classical definition of semaphores with busy waiting.
- If a semaphore value is negative, its magnitude is the number of processes currently waiting on the semaphore.

*-2 = 2 in waiting q*



## Semaphore Implementation with no Busy waiting (Cont.)

```
typedef struct{
    int value;
    struct process *list;
} semaphore;

wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

Noticing that

- Increment and decrement are done before checking the semaphore value, unlike the busy waiting implementation
- The **block()** operation suspends the process that invokes it.
- The **wakeup(P)** operation resumes the execution of a suspended process P.

*← waking up doesn't mean it runs 100%.  
it only means that it gets the chance to run.*





## Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes (to be examined in Chapter 8)
- Let  $s$  and  $Q$  be two semaphores initialized to 1

| $P_0$                    | $P_1$                    |
|--------------------------|--------------------------|
| <code>wait(S) ;</code>   | <code>wait(Q) ;</code>   |
| <code>wait(Q) ;</code>   | <code>wait(S) ;</code>   |
| ...                      | ...                      |
| <code>signal(S) ;</code> | <code>signal(Q) ;</code> |
| <code>signal(Q) ;</code> | <code>signal(S) ;</code> |

- Consider if  $P_0$  executes `wait(S)` and  $P_1$  executes `wait(Q)`. When  $P_0$  executes `wait(Q)`, it must wait until  $P_1$  executes `signal(Q)`. However,  $P_1$  is waiting until  $P_0$  executes `signal(S)`. Since these `signal()` operations will never be executed,  $P_0$  and  $P_1$  are **deadlocked**. This is extremely difficult to debug
- **Starvation** – **Indefinite blocking**
  - A process may never be removed from the semaphore queue, in which it is suspended. For instance, if we remove processes from the queue associated with a semaphore using LIFO (last-in, first-out) order or based on certain priorities.



## End of Chapter 6

